

Medical imaging project - Classifiers

Dr. Veronika Cheplygina

Medical imaging project - timeline

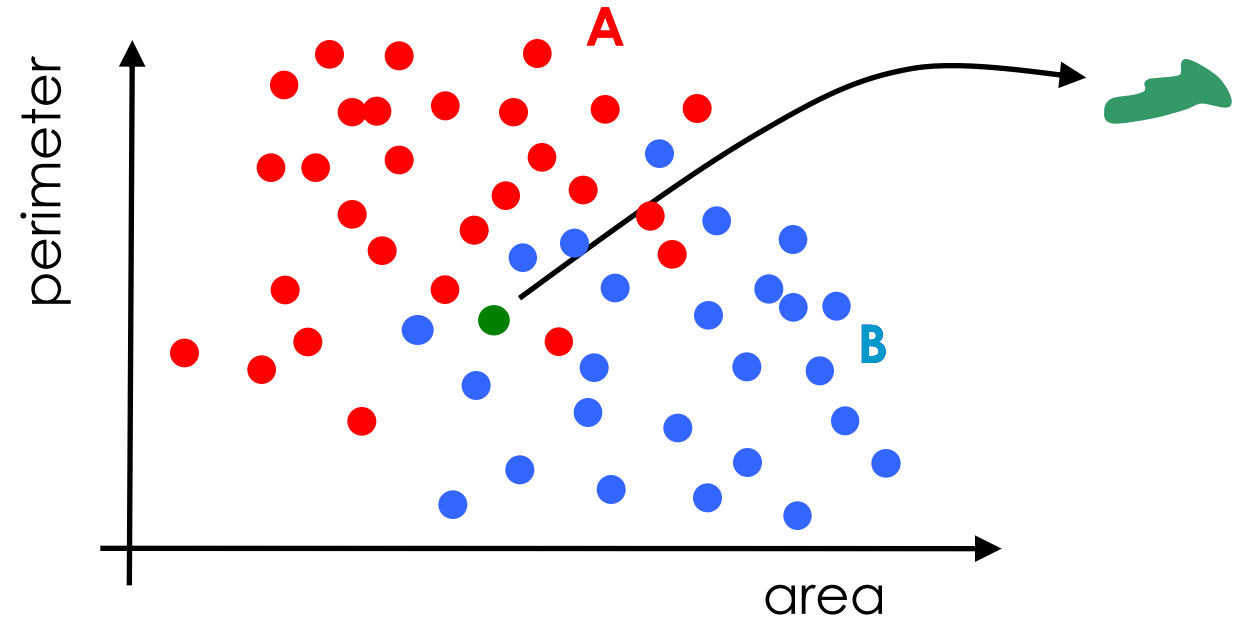


- Tue 6 – Introduction
- Fri 9 – Extract features from images
- Tue 13 – **Classify image to predict diagnosis**
- Fri 16 – Evaluate results
- Tue 20 – TBA / Recent research

Today

- Nearest mean & nearest neighbor classifiers
 - <https://www.youtube.com/watch?v=EK55c9r-n8w>
- Classifier properties
- Overfitting


Nearest mean & nearest neighbor

- Use the features of existing data to train a **classifier**



Example classifier – Nearest mean classifier

- Represent each class by its average point
- Measure distances to “average red” and “average blue”



Example with 5 training points using 3 features

Class 1:

[3 4 4]

[5 5 2]

[1 6 3]

Test point:

[5 3 4]

Class 2:

[2 2 4]


[2 4 6]

Nearest mean classifier - Steps

- Model each class **by its mean**
 - $[3 \ 5 \ 3]$ for class 1, $[2 \ 3 \ 5]$ for class 2
- For the test point, calculate its distances to the class means
 - $\sqrt{9}$ to class 1, $\sqrt{10}$ to class 2
- Find class with the smallest distance
 - Answer is **class 1**

- **Nearest neighbor classifier**

- Model each class by all of its training points (“lazy learning”)
- For a test point, calculate its distances to all training point, find class with the smallest distance



- Example nearest neighbor

Class 1:

[3 4 4]

[5 5 2]

[1 6 3]

Class 2:

[2 2 4]

[2 4 6]


Test point:

[5 3 4]

Class 1


Classifier accuracy / error

- Accuracy = fraction of points that are correctly classified
- Error = $1 - \text{accuracy}$
- [More metrics in next lecture]




Nearest neighbor classifier

- Local changes in training data have large influence on the boundary
- If test data = training data, error = 0




Nearest neighbor classifier


- To reduce variability, look at majority of k nearest neighbors
- Use odd k to avoid ties
- Less sensitive to local changes / boundary is more smooth
- Q: what happens if test data = training data?



- Q: What happens if k is equal to number of training points?




- In theory error will first decrease with k , then increase again
- This is not always the case with real-life data



Nearest neighbor classifier



- **Not** scale invariant
 - E.g. features “weight” and “height”
 - If weight in kg, height in m, height has almost no influence on distance



kNN sensitive to irrelevant features

Take dataset with 2 informative features, add random noise features

Error of kNN increases



Other (not deep learning) classifiers

Linear discriminant

Logistic

Decision tree

Random forest

Support vector machine



...

Classifier properties

“Classifier” can refer to:

- The general type of function you could fit on your data
 - Also “untrained classifier”
- A specific classifier already trained on some data
 - Also “trained classifier”

- Any trained classifier can be visualized in 2D or 3D
 - For each (hypothetical) point, find class it belongs to
 - Draw boundary

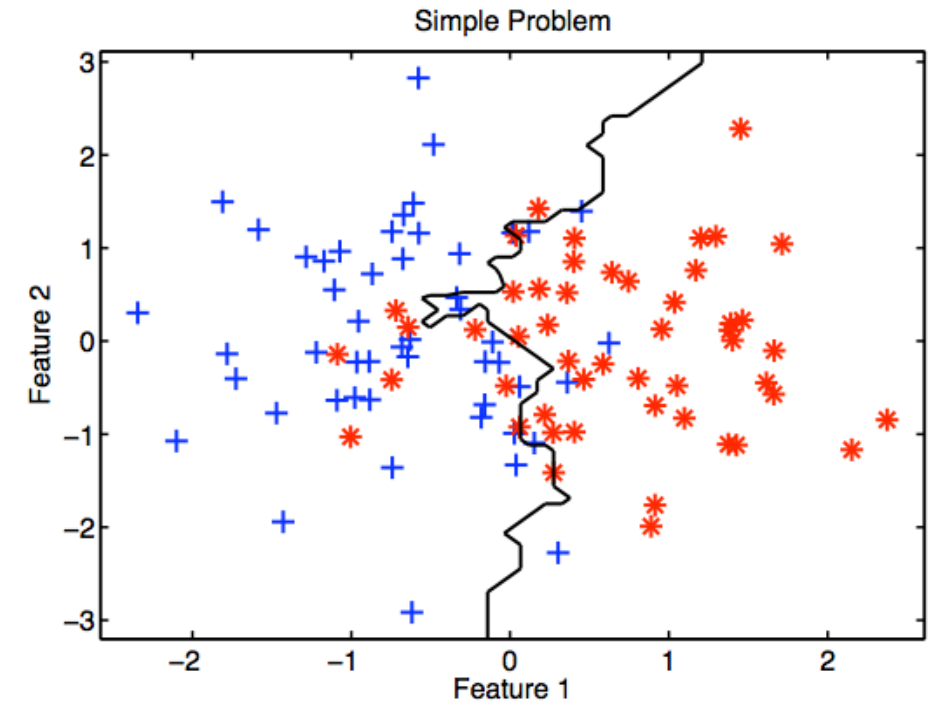



Classifier properties

- Classifiers can be linear, or non-linear
- Some classifiers take all features into account, others they decide the importance of features during training
- Some classifiers have hyperparameters
- Classifier flexibility/complexity

Feature scaling

- Different features might have different scales, for example
 - Color between 1 and 5
 - Area between 1000 and 5000
- Features with larger scales will have more influence on classifier




Feature scaling

- Normalize features to ensure each feature is equally important
- For example, subtract mean, and divide by standard deviation


How to choose a classifier?

- We want to have good separation of the categories (low class overlap)
- It is always possible to design a classifier that perfectly separates your training data
- But, that classifier might **not generalize** to future data, i.e. it **overfits**



How to choose a classifier?

- Use different subsets of data for
 - Training
 - Validation (practice test set)
- Goal: similar performances training & validation sets
 - Bad: 100% training, 50% validation
 - Better: 75% training, 75% validation
- Use these sets to diagnose overfitting (decide on classifier, feature selection etc)





How to choose a classifier?



- When training & validation performances are similar, we expect similar performances on “true” test data
- Use a third subset - test set – to verify you made a good choice
 - Only use for reporting!
 - Do not change method after

Overfitting – 3 reasons


Overfitting

- The classifier has learnt patterns in the training data (training error is low)
- These are not the true patterns / they do not apply to the test data
- The classifier can **not** create good predictions for the test data (test error is high)






Generalization vs overfitting



Overfitting when


- 1) Too few representative training examples
- 2) ...

Overfitting reason 2 – Let's look at this example dataset




We use the “histogram-based classifier” (not used in practice, but related to kNN)

Divide space into cells, label each cell by majority of samples




Number of bins determines how simple/complex the model is


- Underfitting: model is too simple, cannot capture the patterns in the data
- High training error, high test error




We can increase complexity (and fit the data better) by adding more cells



This is a better model, but there are still some errors




Increasing complexity again will fit training data perfectly, but the decisions will be noisy - overfitting has occurred



Overfitting


- More cells = more complexity = can more easily capture patterns in data



Overfitting when


- 1) Too few representative training examples +
- 2) Too complex classifier +
- 3) ...

The complexity the classifier can have, depends on the number of samples



Trade-off sample size and classifier complexity

- Use simple classifiers (e.g. nearest mean or another linear classifier) when data is small




Nearest mean could be the blue classifier, 1-nearest neighbor the red one

Overfitting reason 3 –


What if instead of adding more complexity to the model, we add more features?

Imagine we add a third feature, “how close is this point to the center of the data”

Blue points will have lower values on this feature (closer to center)




- In 2D we need at least 16 samples to fill each cell
- With a third dimension, we need more samples (64 if we use 4 bins)



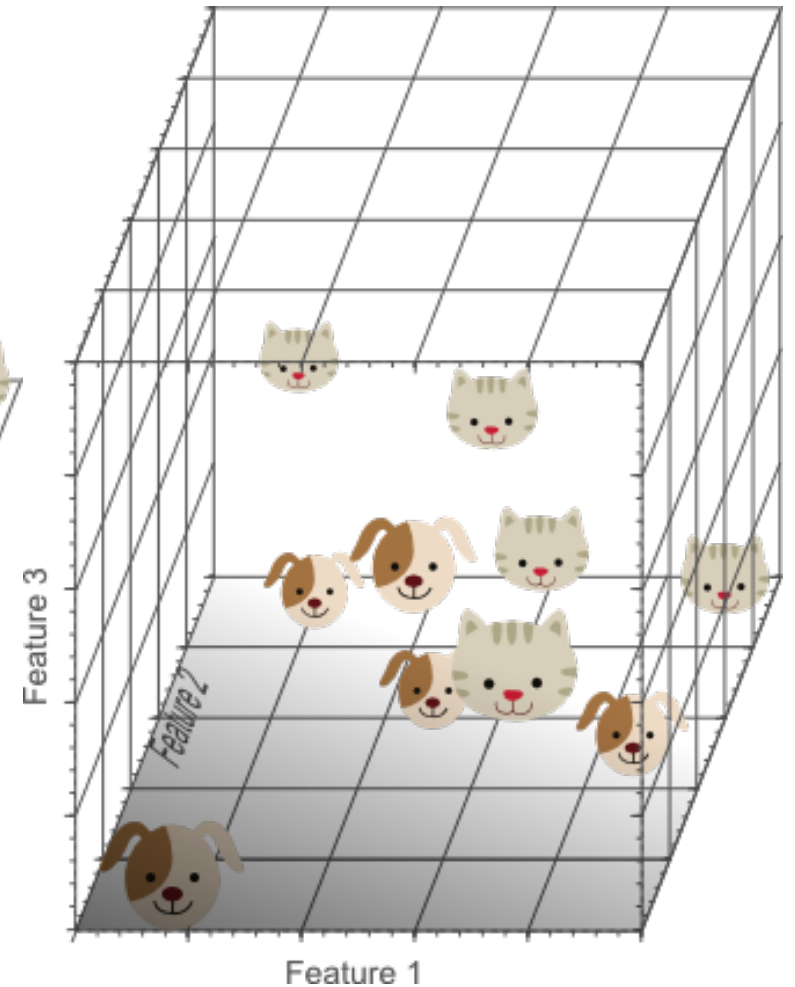


“Curse of dimensionality”

The more dimensions you add, the “emptier” the space becomes, the easier it is to perfectly fit a model to the training data



Further reading:


<https://www.visiondummy.com/2014/04/curse-dimensionality-affect-classification/>




Thought exercise

- Consider 100 uniformly distributed points in a unit line / square / cube ...
- Your goal is to find 1 nearest neighbor (1% of the points, $V = 0.01$)
- How far (what is the value of r) do you need to travel on a line? ($m = 1$)
- How far to travel per dimension, when $m = 2$?

$$r^m = V$$



- In 10 dimensions, need to explore 70% of each dimension to find a neighbor!




Overfitting when

- 1) Too few representative training examples +
- 2) Too complex classifier +
- 3) Too many features


- Generalization
 - Training error \approx test error
- Overfitting
 - Training error \ll test error
 - Too few (representative) objects
 - Too complex classifier
 - Too many dimensions

K-NN & generalization/overfitting:

- Likely overfitting when $k=1$
- Likely underfitting when $k=N$



- Because of its complexity, kNN is more suitable for datasets with many training objects
 - E.g. number recognition
 - Pixel as feature, PCA to 0.9 variance, 5-NN → 1% error




For datasets with few training objects, linear classifiers like **nearest mean** might be more suitable

- Brain data (79 in 2 classes)

Stoeckel, Jonathan, et al. "Classification of SPECT images of normal subjects versus images of Alzheimer's disease patients." *Medical Image Computing and Computer-Assisted Intervention–MICCAI 2001*. Springer Berlin Heidelberg, 2001.

- Mass spectrometry (57 in 3 classes)

Veenman, Cor J., and Annabel Bolck. "A sparse nearest mean classifier for high dimensional multi-class problems." *Pattern Recognition Letters* 32.6 (2011): 854-859.



- We can “diagnose” generalization or overfitting by looking at the training error and test error
- **But** by looking at the test error, the test data becomes part of the training data!
- Use train & validation to choose classifier, test only for reporting
- More on this in the next lecture